

Predicting Fuel Efficiency (Decision Tree Regression and Linear Regression)

In [139]..

```
# Import packages
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.tree import DecisionTreeRegressor
```

Load the data as a Pandas data frame and ensure that it imported correctly.

In [63]:

```
# Load data into dataframe
# Link
link = '/Users/Malloryh5/Downloads/auto-mpg.csv'
# read csv file
data = pd.read_csv(link)
# Print head
data.head()
```

Out[63]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	ford torino

Prep the data for modeling

In [64]:

```
# Drop unneeded columns
data = data.drop('car name', axis=1)
```

Find data type for all data. If strings found replace any strings with the column mean.

In [65]:

```
# Review data types
data.dtypes
```

Out[65]:

```
mpg                float64
cylinders           int64
displacement        float64
horsepower          object
weight              int64
acceleration        float64
model year          int64
origin              int64
dtype: object
```

In [66]:

```
# Review horsepower data
data['horsepower'].unique()
```

Out[66]:

```
array(['130', '165', '150', '140', '198', '220', '215', '225', '190',
       '170', '160', '95', '97', '85', '88', '46', '87', '90', '113',
       '200', '210', '193', '?', '100', '105', '175', '153', '180', '110',
       '72', '86', '70', '76', '65', '69', '60', '80', '54', '208', '155',
       '112', '92', '145', '137', '158', '167', '94', '107', '230', '49',
       '75', '91', '122', '67', '83', '78', '52', '61', '93', '148',
       '129', '96', '71', '98', '115', '53', '81', '79', '120', '152',
       '102', '108', '68', '58', '149', '89', '63', '48', '66', '139',
       '103', '125', '133', '138', '135', '142', '77', '62', '132', '84',
       '64', '74', '116', '82'], dtype=object)
```

In [67]:

```
# Replace any data points that is not a number with nan
data['horsepower'] = pd.to_numeric(data['horsepower'], errors='coerce')
```

In [68]:

```
# Replace any nan in horsepower with mean
data['horsepower'] = data['horsepower'].fillna(data['horsepower'].mean())
```

In [69]:

```
# Create dummy variables for the origin column.
data = pd.get_dummies(data, columns=['origin'])
```

In [70]:

```
data.head()
```

Out[70]:


	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin_1	origin_2	origin_3
0	18.0	8	307.0	130.0	3504	12.0	70	True	False	False
1	15.0	8	350.0	165.0	3693	11.5	70	True	False	False
2	18.0	8	318.0	150.0	3436	11.0	70	True	False	False
3	16.0	8	304.0	150.0	3433	12.0	70	True	False	False
4	17.0	8	302.0	140.0	3449	10.5	70	True	False	False

Create a correlation coefficient matrix and/or visualization. Are there features highly correlated with mpg?

In [71]:


```
# Make a correlation coefficient matrix
corr_matrix = data.corr()
```

```
In [94]: # make a heat map to view the correlation matrix
plt.figure(figsize=(12,12))
# Make heatmap
sns.heatmap(corr_matrix, annot=True, cmap= 'viridis')
# Rotate x-axis labels
plt.xticks(rotation=34)
# Rotate y-axis labels
plt.yticks(rotation= 0)
# Show graph
plt.show()
```



Plot mpg versus weight. Analyze this graph and explain how it relates to the corresponding correlation coefficient.

```
In [103... # Make scatterplot of mpg vs weight
## Figure size
plt.figure(figsize=(10,7))
## Scatter plot
sns.scatterplot(data=data, x='weight', y='mpg', color='blue')
## Title
plt.title('MPG vs. Weight')
## Show results
plt.show()
```



```
In [104... # Analyze te relationship
print(corr_matrix['mpg']['weight'])

-0.8317409332443344
```

There is a negative relationship between the mpg and weight of a car. The more a car weighs the higher the gas milage.

Randomly split the data into 80% training data and 20% test data, where your target is mpg.

```
In [113... # Variable that revomes mpg
X = data.drop('mpg', axis = 1)
# Variable that is only mpg
y = data['mpg']
```

```
In [123... # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state = 0)
```

Train an ordinary linear regression on the training data.

```
In [144... # Save linear regression as variable. Allow linear regression
lin_reg = LinearRegression()
# Fit the train data to reduce RSS
lin_reg.fit(X_train, y_train);
```

Out[144]: ▾ LinearRegression
LinearRegression()

```
In [126... # Make linear regression prediction using train data
y_train_pred = lin_reg.predict(X_train)
# Make linear regression prediction using test data
y_test_pred = lin_reg.predict(X_test)
```

Calculate R2, RMSE, and MAE on both the training and test sets and interpret your results.

```
In [129... # Test Matrix

## Calculate R^2 test
r2_test = r2_score(y_test, y_test_pred)
## Calculate RMSE test
rmse_test = mean_squared_error(y_test, y_test_pred, squared=False)
## Calculate MAE test
mae_test = mean_absolute_error(y_test, y_test_pred)
```

```
In [131... # Train Matrix

## Calculate R^2 test
r2_train = r2_score(y_train, y_train_pred)
## Calculate RMSE test
rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
## Calculate MAE test
mae_train = mean_absolute_error(y_train, y_train_pred)
```

```
In [136... # Print test and training results
print(f'MPG training set: R2={round(r2_train,3)}, RMSE={round(rmse_train,3)}, MAE={round(mae_train,3)}')
print(f'MPG test set: R2={round(r2_test,3)}, RMSE={round(rmse_test,3)}, MAE={round(mae_test,3)}')
```

MPG training set: R2=0.821, RMSE=3.282, MAE=2.48
MPG test set: R2=0.827, RMSE=3.315, MAE=2.7

The model is a good predictor. The R^2 is high, while the MAE and RMSE are low.

1. **R^2 :**
 - Test: .821 or 82.1%
 - Train: .827 or 82.7%
 - 82% of the dependent variable can be explained by the independent variable. This means that the prediction matches 82% of the data.
2. **MAE:**
 - Test: 3.282
 - Train: 3.315
 - On average, the difference of the predictions sum is only 3.3 (test) and 3.2 (train).
3. **RMSE:**
 - Test: 2.48
 - Train: 2.7
 - The predictions differ from the values of only about 2.5.

Pick another regression model and repeat the previous two steps. Note: Do NOT choose logistic regression as it is more like a classification model.

Train an ordinary Decision Tree Regression on the training data.

In [145...

```
# Allow decision tree regressor
tree_regression = DecisionTreeRegressor()
# Fit to data
tree_regression.fit(X_train, y_train);
```

Out[145]:

▼ DecisionTreeRegressor

DecisionTreeRegressor()

In [148...

```
# Make decision tree regressor for test data
y_test_pred_2 = tree_regression.predict(X_test)
# Make decision tree regressor for train data
y_train_pred_2 = tree_regression.predict(X_train)
```

Calculate R^2 , RMSE, and MAE on both the training and test sets and interpret your results.

In [151...

```
## Calculate  $R^2$  test
r2_test_2 = r2_score(y_test, y_test_pred_2)
## Calculate RMSE test
rmse_test_2 = mean_squared_error(y_test, y_test_pred_2)
## Calculate MAE test
mae_test_2 = mean_absolute_error(y_test, y_test_pred_2)
```

In [153...

```
## Calculate  $R^2$  train
r2_train_2 = r2_score(y_train, y_train_pred_2)
## Calculate RMSE train
rmse_train_2 = mean_squared_error(y_train, y_train_pred_2)
## Calculate MAE train
mae_train_2 = mean_absolute_error(y_train, y_train_pred_2)
```

In [155...

```
# Print test and training results
print(f'MPG training set: R2={round(r2_train_2,3)}, RMSE={round(rmse_train_2,3)}, MAE={round(mae_train_2,3)}')
print(f'MPG test set: R2={round(r2_test_2,3)}, RMSE={round(rmse_test_2,3)}, MAE={round(mae_test_2,3)}')
```

MPG training set: R2=1.0, RMSE=0.0, MAE=0.0
MPG test set: R2=0.864, RMSE=8.611, MAE=2.146

The model is overfitted. The train data is higher then the test data. This is not a good regression model to use.

In []: